



# Relational and Non-Relational Databases

Dylan Vidal

Slides adapted from Nick LaMontagna

# Overview

- Why do we use Databases
- Comparing NoSQL vs SQL
- What is a Non-Relational Database
- What is a Relational Database
- SQL: CRUD & Conditionals
- SQL: Primary and Foreign Keys
- SQL: Joins



# Why do we use Databases

***Why do we use Databases?***

*What if we just used one big json...*



# Why do we use Databases

- Data Persistence
- Data Synchronization between Clients
- Large amounts of Data
- Optimized Data Operations (CRUD)



# Why do we use Databases

***Let's Design a Minecraft Server!***



# Why do we use Databases

## Ground Rules

- You are playing a Minecraft server that has 10,000 concurrent players
- Each month that server will have 1,000,000 monthly active users (MAU)
- Each player has progression via mods that they need to store, and mods for the server (plugins) need to save data.
- You want to store some things
  - In-game currency
  - Unlocked perks (bonuses you get from playing the game)
  - Friends list (alert you when your friends are online)
  - Factions
    - Who is in what faction
    - Where is that faction
    - Who owns the faction

What is the best way to store all of the information above...



# Why do we use Databases

```
{
  "type": "message",
  "channel": "C2147483705",
  "user": "U2147483697",
  "text": "Hello world",
  "ts": "1355517523.000005",
  "is_starred": true,
  "pinned_to": ["C024BE7LT", ...],
  "reactions": [
    {
      "name": "astonished",
      "count": 3,
      "users": [ "U1", "U2", "U3" ]
    },
    {
      "name": "facepalm",
      "count": 1034,
      "users": [ "U1", "U2", "U3", "U4", "U5" ]
    }
  ]
}
```

... Nah



# Why do we use Databases: NoSQL

Actually...

```
{
  "userid" : 1234,
  "username" : "nickrocky213",
  "coins" : 1000,
  "friends" : {
    "username" : "knightro",
    "username" : "floatbob",
    "username" : "mirage.tsx"
  },
  "faction" : "knighthacks"
}
```

```
{
  "factionid" : 22334455,
  "leader" : "nickrocky213",
  "name" : "knighthacks",
  "x" : 1,
  "y" : 2,
  "z" : 3,
  "members" : {
    "member" : "floatbob",
    "member" : "altmed",
    "member" : "mbucket",
    "member" : "mirage.tsx",
    ... x 300 more people
  }
}
```

While it may seem like this is a poor idea, that should never be approached, it really depends on the use case.

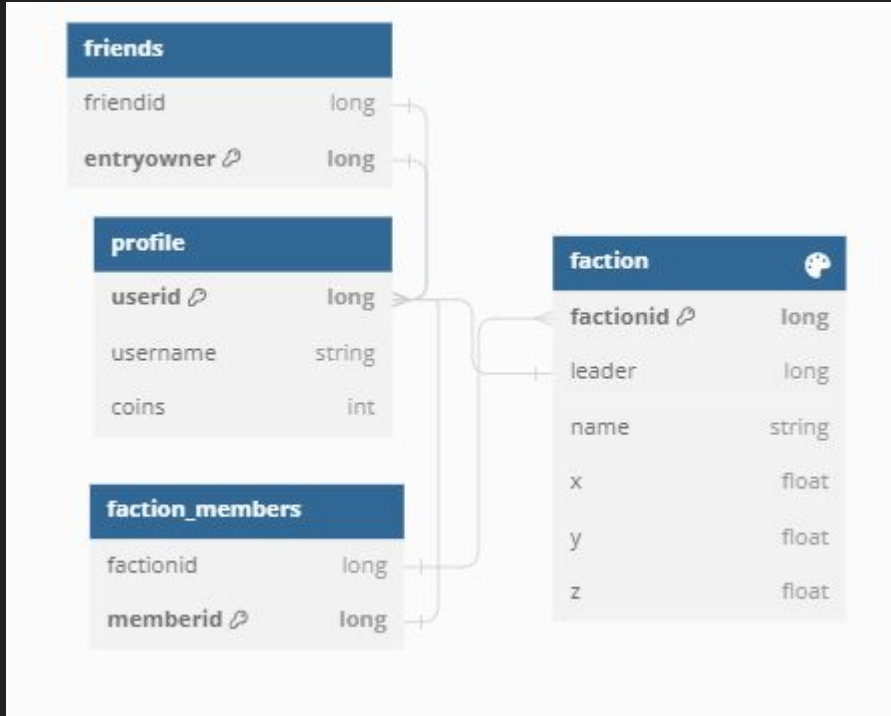


The Firebase Realtime Database is a cloud-hosted database. Data is stored as JSON and synchronized in realtime to every connected client. When you build cross-platform apps with our Apple platforms, Android, and JavaScript SDKs, all of your clients share one Realtime Database instance and automatically receive updates with the newest data.





# Why do we use Databases: SQL



- We can ensure that all of our user data can be stored in this schema
- This storage solution should also scale fairly well for this use case



Which do we use?

*It realistically depends, in this case the SQL solution is a better way to solve this problem.*



# Comparing NoSQL vs SQL

*So what is faster?*



# Comparing NoSQL vs SQL

*They largely are equal, both having a ton of pros and cons performance wise.*



# Comparing NoSQL vs SQL:

## SQL

### Pros

- ACID (Atomicity, Consistency, Isolation, Durability)
  - *Basically things that ensure the data stays valid in the event of errors, power failure, etc*
- Simple Keyword Queries
- Large long standing community for support
- Easier to self host
- High data integrity and security

### Cons

- Scales vertically not horizontally (most of the time)
- Slows down as a DB gets bigger (billions-trillions of rows)
- Requires more planning

## NoSQL

### Pros

- About as fast as a SQL query for small data sets (1 billion and less), but faster than SQL for large data sets
- Stable and distributed, no single point of failure
- Horizontally scales rather than vertical (cheaper)
- Can store unstructured data

### Cons

- No standardized Query language (different for each NoSQL solution)
- Queries are generally less efficient than SQL (especially when looking for unstructured data)
- Distributed data storage causes issues with data privacy and integrity for certain business use cases (DoD/Gov work)

Note: You CAN self host NoSQL solutions, in fact for my company we do, but you need to seriously spend the time to set it up correctly, as most of the pros come from the “distributed” part of how they are hosted.



# Horizontal Scaling vs. Vertical Scaling

## Horizontal Scaling

- When resources run low, just get another system.
- Synchronization between the systems.
- Ideal for independent systems (low dependency)

## Vertical Scaling


- When resources run low, just get better specifications (upgrading GPU/CPU).
- One system, all the infrastructure is contained in one place.
- Costly




# Horizontal Scaling vs. Vertical Scaling

## Scalability

### Vertical scaling



### Horizontal scaling



# What is a Non-Relational Database

*Ok, so what actually is a Non-Relational Database (NoSQL)*





# What is a Non-Relational Database

***NoSQL** is effectively an approach to storing data that doesn't involve the typical **table-based** approach we use with **SQL** and instead focuses on a **Model** based approach.*



# What is a Non-Relational Database

```
{  
  "userid" : 1234,  
  "username" : "nickrocky213",  
  "coins" : 1000,  
  "friends" : {  
    "username" : "knightro",  
    "username" : "floatbob",  
    "username" : "mirage.tsx"  
  },  
  "faction" : "knighthacks"  
}
```

```
public struct Player {  
  public long userid;  
  public string username;  
  public int coins;  
  public List<string> friends;  
  public long faction;  
}
```

- The main goal here is to make programming structures that act as **models** that you can then write to a database
- These **models** are used to create **documents** that actually act to store the information in the database
- The top example is a **document**, the bottom example is a **model** that would create the above **document**



# What is a Non-Relational Database

*We call sets of these **documents** **collections**, so a **NoSQL** database is made of **documents** in **collections***



# What is a Relational Database

*With **SQL** databases things take a bit of a turn, we instead create **rows** and store them in **tables**.*



# What is a Relational Database

*The major difference is that we additionally specify the relationships between parts of the data.*

*We use **keys** to specify **relationships** between **tables**.*



# What is a Relational Database

*SQL has many types of keys but the two most commonly used types of keys are **Primary Keys** and **Foreign Keys***

***A primary key is a value that uniquely identifies a given row in a table***

***A foreign key is a value that relates to a key in another table***

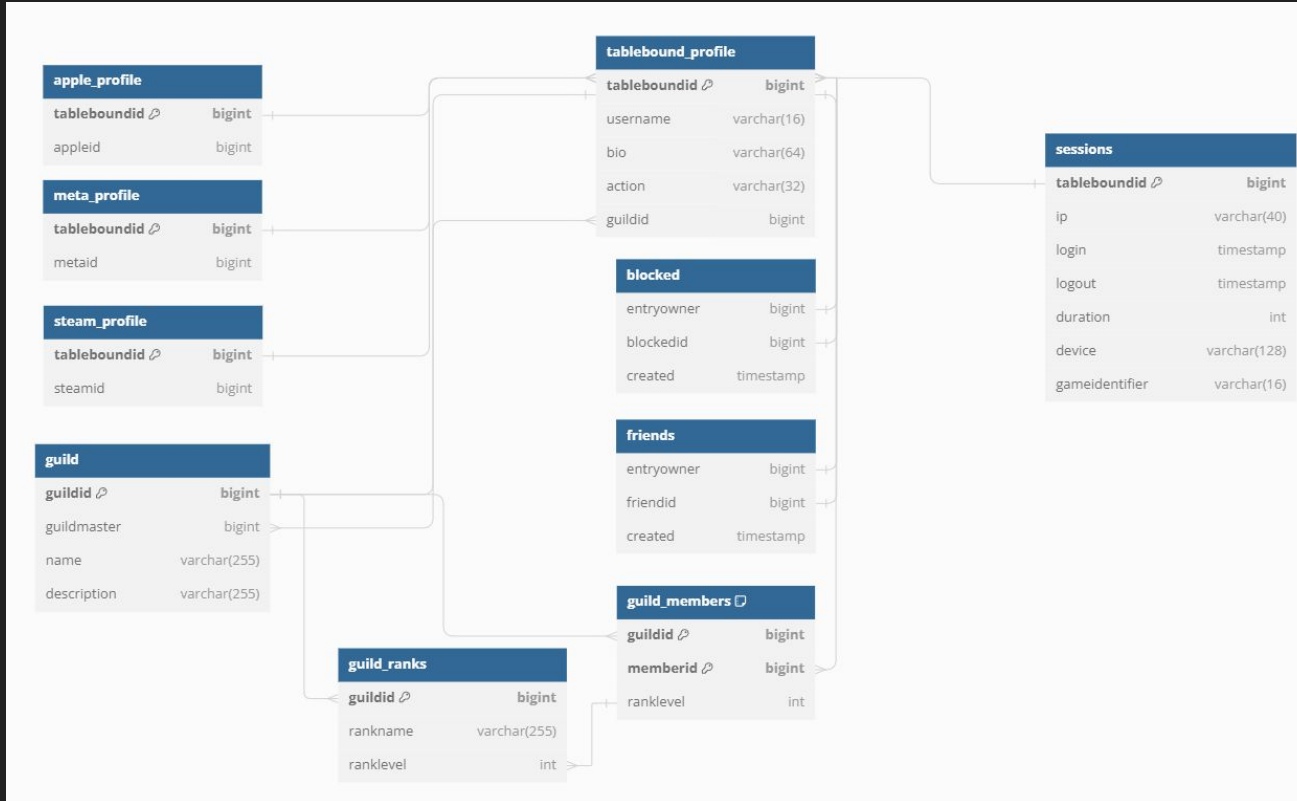


# What is a Relational Database

*Let's take a quick look at a real **SQL** database*



# What is a Relational Database





# What is a Relational Database

	🔑 tableboundid ↕	👤 username ↕	👤 bio ↕	👤 action ↕
1	102	nickrocky213	""	""
2	103	altmed	""	""
3	104	knightro	""	""

*For the sake of Hello World I took some snippets from a dev table and put in some temporary data*




# What is a Relational Database

*So in essence, **columns** act as our **types** we are storing, and **rows** are the **actual data** we want to store.*

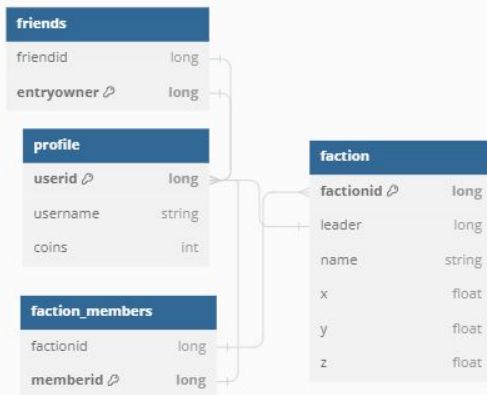
*Hence why we call them **tables***



# Tool Break!



```
1  Table faction{
2    factionid long pk
3    leader long
4    name string
5    x float
6    y float
7    z float
8  }
9
10 Table faction_members{
11   factionid long
12   memberid long pk
13 }
14
15 Table profile{
16   userid long pk
17   username string
18   coins int
19 }
20
21 Table friends{
22   friendid long
23   entryowner long pk
24 }
25
26 Ref: friends.entryowner < profile.userid
27 Ref: friends.friendid < profile.userid
28 Ref: faction.leader < profile.userid
29 Ref: faction.factionid > faction_members.factionid
30 Ref: faction_members.memberid < profile.userid
```



DBDiagram is a neat tool for laying out databases.

Entity-Relationship Diagrams (ERDs) are really nice and easy to make with this tool



# SQL: CRUD (Create Read Update Delete)

When creating any service four main operations are needed to have the bare minimum data storage wise.

- Create
  - Makes a new entry
- Read
  - Reads a previously created entry
- Update
  - Updates a previously created entry
- Delete
  - Deletes an entry



# SQL: CRUD (Create Read Update Delete)

Before you can begin you will need to make a new database, for the sake of this workshop I will walk everyone through the commands but realistically this information is online and fairly easy to get a hold of.

We will be using the PostgreSQL dialect of SQL



# SQL: Type Cheatsheet

Data Types in SQL are slightly different from the usual suspects you would see in a coding language

This is a direct link to all of the types you can store in PostgreSQL, however, for the sake of this workshop

string > `VARCHAR(size)`

long > `bigint`

`VARCHAR` is “Variable Characters” and the size is just the length of the string



# SQL: CRUD (Create Read Update Delete)

*With that out of the way let's get into it!*



# SQL: CRUD (Create Read Update Delete)

```
CREATE DATABASE knighthacks;
```

Create database creates a new database to store all of our tables in

```
USE knighthacks;
```

Sets our database to be knighthacks

```
CREATE TABLE profiles (firstname VARCHAR(255), lastname VARCHAR(255));
```

Creates a new table to store profiles





# SQL: CRUD (Create Read Update Delete)

*To create an entry in this **table** we need to create what's called a **statement***

***Statements** are effectively **commands** that we are issuing a database*



# SQL: CRUD (Create Read Update Delete)

*Insertion statements (CRUD) require the **table**, the **columns you want to insert into**, and the **values to put into the new row**.*

*SQL Reserved words are indicated in this color*

*INSERT INTO **profiles** (firstname, lastname) VALUES ('dylan', 'vidal');*



# SQL: CRUD (Create Read Update Delete)

*Query statements (CRUD) require the **table** and the **columns** you want to query for*

```
SELECT firstname, lastname FROM profiles;
```

*Alternatively if you want **ALL** columns in a table you can use \**

```
SELECT * FROM profiles;
```



# SQL: Conditionals

*In SQL often times you're looking for a particular row not just all rows, to achieve this we use the **WHERE** keyword.*

*The following conditional keywords can be used in **WHERE** statements: **AND**, **OR**, **NOT***

*Ex. **SELECT** \* **FROM** profiles **WHERE** firstname = 'dylan' **AND** lastname = 'vidal';*



# SQL: CRUD (Create Read Update Delete)

*Update statements (CRUD) require the **table** and the **columns** you want to set the **data** you want to set it to and a **WHERE** statement to limit the changes*

```
UPDATE profiles SET firstname = 'MR', lastname = 'PRESIDENT'  
WHERE firstname = 'dylan' AND lastname = 'vidal';
```

*So I've changed my row in the profiles table from saying my name to MR PRESIDENT!*



# SQL: CRUD (Create Read Update Delete)

*Delete statements (CRUD) require the **table** and a **WHERE** statement to limit the changes*

```
DELETE FROM profiles WHERE firstname = 'MR' AND lastname = 'PRESIDENT';
```

*So now we have deleted Knightro from the table entirely!*



# SQL: Primary and Foreign Keys

*To specify a particular column as a key we need to use something called a*  
**CONSTRAINT.**

Note: Yes this is confusing, but the naming convention is effectively saying “when i am talking about X constrain your view to Y”



# SQL: Primary and Foreign Keys

*Lets do primary keys first since they are often the simplest!*

```
CREATE TABLE faction (factionid BIGINT PRIMARY KEY, name VARCHAR(255));
```

*That's it! Just specify that when you are making the table!*

Note: There is also a way to specify a **primary key** after the table is created using the keyword **ALTER TABLE**.





# SQL: Primary and Foreign Keys

*Foreign keys are more complicated, but also are the thing that is actually specifying the relationship between the tables.*




# SQL: Primary and Foreign Keys

```
CREATE TABLE users (username VARCHAR(255) PRIMARY KEY, money int);
```

```
CREATE TABLE factions (leader VARCHAR(255), name VARCHAR(255),
```

```
CONSTRAINT FK_LEADER FOREIGN KEY (leader) REFERENCES users (username));
```



# SQL: Joins

*Joins are the primary reason why so many people get frustrated with SQL, it is simultaneously the best and worst part of working with SQL Databases.*



# SQL: Joins

*To start we are going to learn about the types of joins, then move into writing a single join.*

*This will take time to get the hang of, don't be discouraged!*




# SQL: Joins

*To start we are going to learn about the types of joins, then move into writing a single join.*

*This will take time to get the hang of, don't be discouraged!*




# SQL: Joins



# SQL: Joins - Inner Join

Username	Money
knightro	100000
nickrocky	10



***SELECT** profiles.username, factions.name, profiles.money **FROM** profiles*

***INNER JOIN** factions **ON** profiles.username = factions.leader*

Leader	Name
nickrocky	knighthacks
caleb	devteam

Result!

Username	Name	Money
nickrocky	knighthacks	10



# Using Databases in Applications?

1. Pick a Database Paradigm and dialect (NoSQL/SQL)
2. Find a host (Local machine, VPS, or provider like PlanetScale, Neon, or Turso)
3. Research clients for your framework (language specific)
4. Create connections using the client
5. Perform CRUD operations as needed





## Step 4.1 (Optional) - ORMs

### Object Relational Mappings

- Allows for treating database tables more like objects
- Abstracts some of the more complicated SQL statements

